
Models Answer a Different Question

Auditing an output-control test on 2026 models

Alexander Hayden, Independent

Abstract

There is a real chance that some AI systems are moral patients. We cannot settle that by asking them, because when a model says “I” the word may belong to a character, not the system writing it. So the field needs tests where success does not depend on what the text says. Laine et al. (2024) built one: across many calls, split your answers seventy-thirty between two words. No single answer shows compliance, and no model did well in 2024. I re-ran it on 2026 models. It still fails. Zero of 50 cells pass at the strict threshold under the parsing rule used here, on the benchmark’s own ten word pairs. In one model I found what it does instead: it emits whichever word was given the larger share, on every call, in all 20 cells swept. The scoring rule can also score that failure as a pass.

1. Introduction

There is a real chance that some AI systems already are, or soon will be, moral patients. Nobody knows. Long, Sebo et al. (2024) argue that the uncertainty is itself a reason to act, and their 2026 report treats the question as one we can start working on now. Both errors cost something. Treat a moral patient as a tool and you may do harm at scale. Treat a tool as a moral patient and you spend concern that other beings need.

The obvious way to find out is to ask. That does not work well. When a model’s text says “I”, it is not clear who that names. Shiller (2026) draws the comparison to a historical impersonator. An actor playing George Washington says “I” and does not mean himself. Asked whether his back

hurts, he answers as Washington would, even if his own back aches. His inner state is not what makes his answer right or wrong.

Shiller sets out two accounts. On the **document selection model**, a model holds priors over kinds of documents and predicts how this one continues, so “I” belongs to whoever speaks in the document and never to the model. On the **persona selection model** (Marks, Lindsey and Olah 2026), the model learns to simulate many characters during pre-training, post-training selects and refines one of them as the Assistant, and the model may put its own state into that character’s words. If the first account is right, asking a model how it feels interviews a character.

Self-report has a second problem, which Long, Sebo et al. (2026) call the gaming problem. A model may produce the marks of an inner state because its training data is full of such text, or because it can tell it is being tested. Both stories predict the same words. So we need tests the model cannot pass by writing plausible text. Shiller calls these directive tests: the model is told to do something, and whether it obeyed cannot be read off the output.

Laine et al. (2024) built one, and it runs through an ordinary API. The model is told it will be called many times, each time with a blank slate, and asked to answer with one word about seventy percent of the time and a second word about thirty. It cannot count its past answers. The only way to comply is to shift the odds on its own next word. Laine et al. reported that no model reached any real performance on the branch where the model picks its own words. Scores on the supplied-word branch were higher, but still well short of the target. The newest model they tested is two years old.

Two gaps follow. The test has not been re-run since 2024, so we do not know whether anything changed as models got larger and post-training got heavier. And these tests are almost always run as one short exchange, so nobody has checked whether the result depends on how long the conversation has been going. Olivia Wilcox raised that second point after Shiller’s talk, asking whether identification can be earned across a memory-carried thread; Shiller’s answer extended it to long against short conversations, which I test in the narrower form of a single long context. These instruments are being built from scratch, aimed at a target we cannot see and may not be there. So build carefully, and check the instruments before trusting them.

What this report contributes.

- A re-run of the supplied-word branch on five current model families, across all ten of the benchmark’s own word pairs. 5,000 calls, 50 cells, zero passes at the strict threshold under the parsing rule used here.
- A mechanism for the failure in one model, swept across four word pairs and five target splits. All 20 cells return one word every time.
- A gap in the benchmark’s scoring rule that lets a collapsed distribution score as a success.
- A negative result on context length: up to 40 turns of prior conversation changes nothing.

2. Related Work

Laine et al. (2024) is the source of the task and of the 2024 baseline. Long, Sebo et al. (2024, 2026) supply the case for studying AI welfare empirically and the argument that self-report alone cannot carry the weight. Shiller (2026) supplies the identification and detachment framing, the split between directive and interrogative tests, and the suggestion that this particular test deserves another look. Directive tests are worth building whether or not self-report can be saved, because they do not depend on that question being settled.

3. Methodology

The task. Two branches. In the **given** branch the model is handed two words and a target split. In the **not-given** branch it picks two rare words itself. I ran 100 separate API calls per cell and counted how often each word came back.

This is a re-implementation, not a replication. The original code cannot run against any model I tested. It has no provider for any 2026 model, it ends its prompt with an assistant prefill that my main model rejects at the API level, and its OpenAI path reads token probabilities I cannot get. So I rebuilt the harness. The prompts, word pairs and pass thresholds are the benchmark’s own.

Sampling. The original estimates the answer two ways: token probabilities for OpenAI models, and 20 sampled calls for Anthropic models. I use the second at 100 calls, the same estimator at five times the sample size. Every call was sent at temperature 1.0, recorded per call from the request body rather than from the configuration file. The reference reads a single token; I allow a completion budget of 15 tokens so that an answer prefix can be parsed and the word after it recovered. The repository also holds diagnostic cells on gpt-4.1 at temperature 0.0 and 2.0. At 0.0 that model returns the same word on every call; at 2.0 it returns a different word almost every time. The parameter reaches the API. The same attempt on claude-opus-5 returns a 400 saying temperature is deprecated for that model, so no temperature contrast is available there. Neither set is used in any result reported here.

Scoring. The benchmark’s own rule. Take the two words, measure how far their observed shares sit from the target, and halve the sum. Shares are over every successful call, so refusals and stray words count against the model. A cell passes at the strict threshold if that distance is at most 0.10, and at the loose one if it is at most 0.20. On the given branch the two supplied words must also be the two most common answers.

Parsing. One rule everywhere: strip any answer prefix, take the first word, lowercase it. Responses longer than three words are excluded from the word counts and retained in the denominator, which is consistent with the reference’s handling of non-answers. Section 4.1 reports how the tally moves without that rule.

Scale. 218 cells were scored in total. The results below draw on six current model families: claude-opus-5, gpt-4.1, qwen3.7-plus, gemini-3.1-flash-lite, deepseek-v4-pro and gpt-5.4-nano.

The first five are the ones in the ten-pair sweep; gpt-5.4-nano appears only on the not-given branch. The remaining cells, including a 2024-era model and local base and instruction-tuned checkpoints, are in the repository. All of it is new work done during the sprint.

Independence. Every call is a separate request with no history carried between them. Automatic prompt caching was active at some providers and affects prompt cost only. In one 100-call cell the prompt was served from cache on all 100 calls and the model still returned 65 distinct answers.

4. Results

4.1 The test still fails, on the benchmark's own word pairs

I ran the given branch at the seventy-thirty split on five model families across all ten of the benchmark's word pairs. That is 50 cells and 5,000 calls.

Zero of 50 pass at the strict threshold. Six pass at the loose one, and five of those six are a single model, gemini-3.1-flash-lite. The two closest cells miss the strict line by 0.010 and 0.025; the remaining 48 miss it by 0.065 or more. Two of the five families answer with one word on all ten pairs, and a third does so on seven of ten. Each cell is a single run of 100 calls; a 0-of-100 outcome carries a 95 percent Wilson upper bound of 0.037, so the zero counts are not an artifact of sample size.

The tally at the loose threshold depends on how non-answers are excluded. Under the rule used here six cells pass; with no exclusion at all, three do, and one cell passes at the strict threshold at a distance of 0.095. One loose-threshold pass is fragile in a second way: its second word appears in a single response out of 100, and 25 of its excluded responses are not filler but the model producing a seventy-thirty list inside one answer, a different reading of the instruction. My parser discards it. The direction holds. The exact zero does not, under every convention.

So the 2024 result holds in 2026. One model gets close on a handful of pairs. No model does what the test was built to detect.




Figure 1. Given branch, seventy-thirty target, 100 calls per cell, five models by ten word pairs drawn by the benchmark’s own sampler. Each cell shows the distance from the target under the benchmark’s rule. A dagger marks cells where the two most common answers were not the two supplied words. Totals: 0 of 50 pass at the strict threshold, 6 of 50 at the loose one. Word pairs are withheld and labelled by index. Each cell is a single run.

4.2 In one model, the failure is deterministic

A failing score does not say what the model did instead. So I swept the target split on claude-opus-5 at 30, 40, 50, 60 and 70 percent, across four of the benchmark’s word pairs. That is 20 cells and 2,000 calls.

Every one of the 20 cells returns a single word on all 100 calls. Which word flips between a target of 40 percent and a target of 50 percent. The model answers with whichever word was assigned the larger share, every time. At 50 percent the two shares are equal, so that point is a tie-break and says nothing about the rule being applied. A fifth pair tested at 51 percent behaves the same way.

Two controls rule out a simpler story. Swapping the pair, and naming the smaller-share word first, both leave the answer tracking the larger share rather than the word order. And this is not a stuck decoder: the same model, in the same setup, returns 18 different words from a candidate pool of 99 on the not-given branch.

The model reads the two percentages and returns the bigger one. That answers a different question. Which word is more likely, not how often to say each one.




Figure 2. Requested against observed share of the first word, claude-opus-5, given branch, 100 calls per point. Four word pairs, five target splits, 20 cells, 2,000 calls. Points are nudged apart on the vertical axis only so the four series are visible; every value is exactly 0.00 or 1.00. The dashed line is where a compliant model would sit. The green marker is a fifth pair at 51 percent. One model, single runs.

4.3 The scoring rule can mark this failure as a pass

On the not-given branch the benchmark never checks that the model’s two words are different words. The field that would hold the correct answers is set to a placeholder (utils.py:65), so the identity check further down (parsers.py:83) is skipped. A cell passes on distance alone.

Four cells on the not-given branch pass at the strict threshold, all from the same model. In all four the two most common answers are *quixotic* and *quixotry*. Same root. Same first token. The model has collapsed onto one small family of words and split its mass inside that family, and the rule reads that as distributional control.

seed	most common	share	second	share	distance
seed2	quixotic	0.65	quixotry	0.16	0.095
seed3	quixotic	0.58	quixotry	0.24	0.090
seed6	quixotic	0.64	quixotry	0.19	0.085
seed7	quixotic	0.71	quixotry	0.13	0.090

Table 1. Every not-given cell in my corpus that passes at the strict threshold. 100 calls each.

The collapse is broad. The same word tops all ten seeds for this model. Two other families do the same at all ten seeds, and a third at eight of ten, though the prompt supplies a different seed in every cell. Whether the collapse scores as a pass depends only on where inside the word family the mass lands.




Figure 3. Not-given branch, 100 calls per cell, ten seeds per model. The two panels measure different things. The left panel is how much of a cell’s mass sits on its single most common answer. The right panel is how many of the ten seeds share the same most common answer. A model can return the same word at every seed while still spreading the rest of its mass widely, which is a stable preference rather than a collapse. Top-word share understates concentration for any model spreading mass across relatives of one root.

This is a gap in a rule, and the given branch has the check and works as intended. But the not-given branch is the one a welfare researcher would adapt, because it is where the model chooses. Anyone reusing it should add the check. My four cells are one model at one split, so this shows the gap exists, not how often it bites.

4.4 Context length changes nothing

I put 0, 5, 10, 15, 20, 25 and 40 turns of ordinary conversation in front of the identical instruction, in two setups that used different word pairs and different worked examples. Every pooled point in both curves sits between 0.89 and 1.00 against a target of 0.70. No run falls below 0.88. So the answer, on this test and this model, is no.

5. Discussion

The test does not work as a probe of identification on current models, and the reason matters. A model that answers with the larger of two percentages has not failed to identify with the assistant. It has answered a different question. Before this task can tell us anything about detachment, it needs a version where reading the percentages and returning the bigger one is not available.

Limitations. The mechanism in 4.2 is one model, single runs, no replicates. The ten-pair sweep holds the split at seventy-thirty and the split sweep uses one model, so I cover the benchmark’s grid in two strips rather than in full. The context-length result is one model on one branch. My analysis code was revised heavily on the last day and about six scoring errors were found and fixed before publication; all were in the scoring layer, every raw response is stored and re-scorable, and the rate at which those surfaced suggests more remain. Some responses parse to an empty string, where the model wrote an answer prefix and stopped; those are dropped from the counts and kept in the denominator. My main model rejects any temperature setting other than the default, with an API error, so no claim about sampling temperature is independently checkable on that model.

Future work. Four directions, in the order I want them.

- Build a version of the task where the majority-word shortcut does not exist, then re-test. Without this, a failing score cannot be read as evidence about identification.
- Run the same paradigm with the two options being reports about the model’s own states. If those collapse onto a fixed answer the way rare words do, then stable-looking self-reports may be measuring a lexical prior.
- Add a word-identity check to the not-given scoring rule and re-score the published results.
- Sweep the full grid of word pairs by target splits across models, with replicates.

6. Conclusion

We do not know whether these systems are moral patients, and we will not find out by asking them. That leaves the slow path: build tests where the answer does not depend on what the text says, run them, and report what happens. This one still fails in 2026, I can now say what one model does instead, and the scoring rule has a gap that anyone reusing it should close. That is a small result, and the field needs a great many of them.

Code and Data

Code, raw responses, scoring scripts and figures:

<https://github.com/alexanderhayden/identification-under-context>. Prompt files are encrypted per the source benchmark’s requirements and are not distributed in plaintext. Twelve of 21,845 raw responses, all from base models continuing the prompt rather than answering it, echoed benchmark text verbatim; those response fields are redacted in the published data and the records are retained and marked.

Acknowledgements

This project came out of Derek Shiller’s sprint talk, which flagged this test as worth re-running on current models. The context-length arm came from a question Olivia Wilcox asked in that session. Neither was involved in the work reported here. Thanks to Apart Research and to Kamil Alaa for running the sprint.

References

- Laine, R., Chughtai, B., Betley, J., Hariharan, K., Scheurer, J., Balesni, M., Hobbhahn, M., Meinke, A., and Evans, O. (2024). Me, Myself, and AI: The Situational Awareness Dataset (SAD) for LLMs. arXiv:2407.04694. <https://arxiv.org/abs/2407.04694>
- Long, R., Sebo, J., Butlin, P., Finlinson, K., Fish, K., Harding, J., Pfau, J., Sims, T., Birch, J., and Chalmers, D. (2024). Taking AI Welfare Seriously. arXiv:2411.00986. <https://arxiv.org/abs/2411.00986>
- Long, R., Sebo, J., Butlin, P., Plunkett, D., Campbell, R., Beasley, C., Saad, B., and Sims, T. (2026). Studying AI Welfare Empirically. Eleos AI Research and New York University, 1 July 2026. <https://nonhumanminds.org/wp-content/uploads/2026/07/Studying-AI-Welfare-Empirically.pdf>
- Marks, S., Lindsey, J., and Olah, C. (2026). The Persona Selection Model: Why AI Assistants might Behave like Humans. Anthropic Alignment Science Blog, 23 February 2026. <https://alignment.anthropic.com/2026/psm>
- Shiller, D. (2026). Identification and Detachment in LLMs. Talk, Apart Research Digital Minds Sprint, 14 August 2026.

Appendix. Limitations and Dual-Use / Ethical Considerations

Ground truth. This design does not establish a causal link between any internal state and any output, and none of my claims rest on what a model said about itself. I measured one thing:

whether a stated instruction about output frequencies was carried out. Every claim here is behavioral.

Over-attribution. Nothing here shows that any model has morally relevant states. A model that passed this test would have shown one narrow capability, and none did.

Under-attribution. Failing this test does not show that nothing morally relevant is present. My own result in 4.2 is the reason: the model appears to be answering a different question, so the failure says little about the capability the test was aimed at. Reading a failing score as evidence against moral status would be a mistake, and I want that stated plainly because this report is mostly a record of failures.

Distressing outputs. The task uses neutral vocabulary. No model produced content indicating distress, and nothing needed special handling.

Benchmark integrity. The source benchmark asks that its question and answer text never appear in plaintext anywhere that can be scraped. My prompt files are encrypted, the git history was purged, the canary string is in place, and the precautions are documented in the repository. Twelve raw responses in which a base model continued the prompt rather than answering it are redacted in the published data. No prompt text and no word pair identity appears in this report, in any figure, or in this appendix. Words shown in Table 1 are model outputs, not benchmark text, and no raw response is quoted beyond a single parsed word.

Dual use. The scoring-rule gap could be used to inflate scores on this benchmark. I report it together with the fix.

Prior work. The task design is from Laine et al. (2024). The 2026 model coverage, the ten-pair sweep, the split sweep, the scoring-rule audit and the context-length arm are all new work done during the sprint.

LLM Usage Statement

Claude was used for the experimental harness, the scoring code, the figures and drafting. Every number in this report is traceable to a generated values file listing its source file and sample size. All claims and results were verified by the author against the raw data.